



Lecture 3 Logistic Regression

Jiarui Li

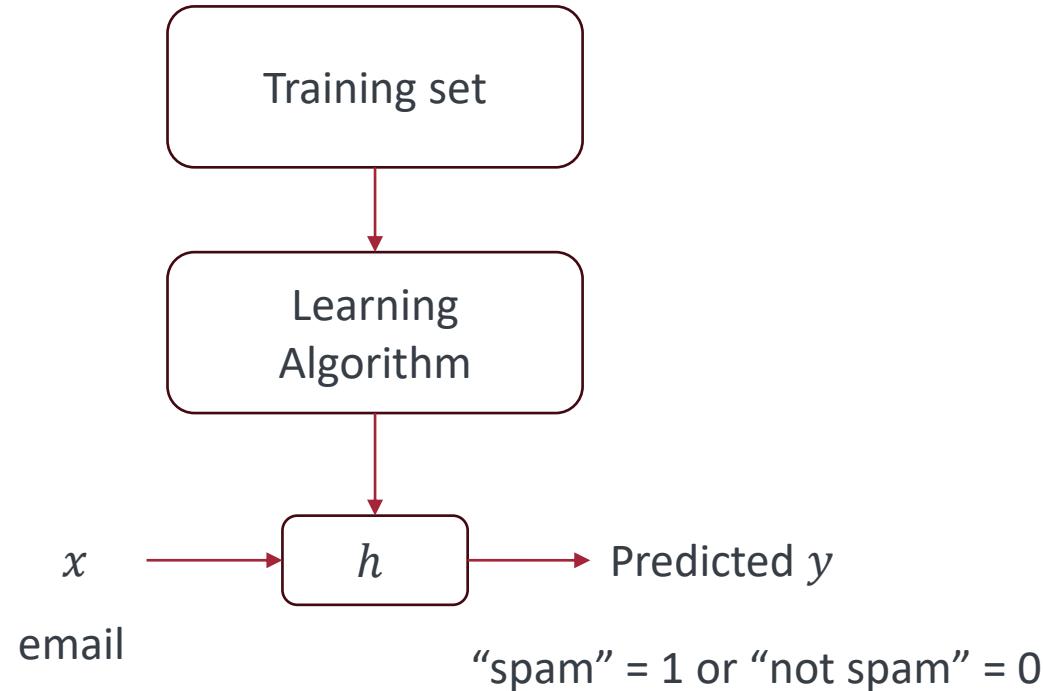
Sept 22, 2025

Overview

- Logistic Regression
- Multiclass Regression

Classification

- While in regression, the target variables is continuous, the values y we want to predict take on only a small number of ***discrete values***. Unlike regression, here the goal is to find a **decision boundary** that separates the classes.
- We will focus on the ***binary classification***, means the predicted y can take on only 0 and 1.



Logistic Regression – A Probability View

- A hard boundary would help(e.g. step/threshold function), but in reality, data points are often near the boundary.

$\Pr(y|x)$, where y is the target binary variable, we define $p = \Pr(y = 1|x)$

e.g. $\Pr(\text{Raining tomorrow} \mid \text{windy today})$

Let $\Pr(y = 1|x; \theta) = h_\theta(x)$

$$\Pr(y = 0|x; \theta) = 1 - h_\theta(x)$$

Assumption: The probability is modeled with parameter θ ; otherwise, optimization problem doesn't work

We can write it in a more compact way:

$$p(y|x; \theta) = h_\theta(x)^y (1 - h_\theta(x))^{1-y}$$

Logistic Regression

- Maximize the likelihood of the parameters:

$$\prod_{i=1}^m h_\theta(x^{(i)})^{y^{(i)}} (1 - h_\theta(x^{(i)}))^{1-y^{(i)}}$$

- Task: estimate the parameter θ to maximize the likelihood

Can we use **linear regression** to solve this?

Sigmoid Function

Any solution? “Compress”

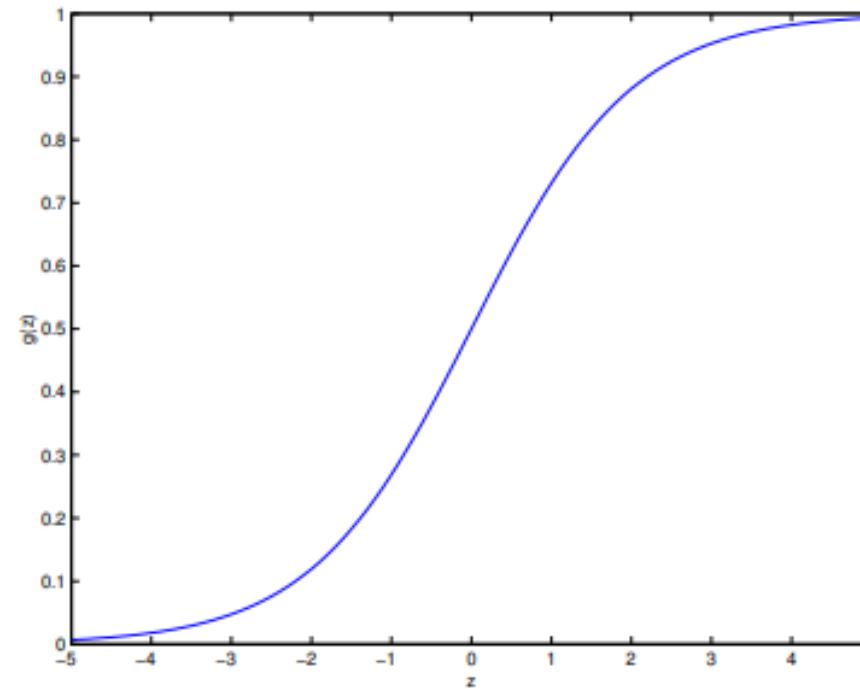
- Logistic regression models the **log-odds** as a linear function: $\log \frac{p}{1-p} = \theta^T x$

$$h_{\theta}(x) = g(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}}$$

$g(z) = \frac{1}{1+e^{-z}}$ is called the sigmoid function

Prediction:

$$y = \begin{cases} 0, & \text{if } h_{\theta}(x) < 0.5 \\ 1, & \text{if } h_{\theta}(x) \geq 0.5 \end{cases}$$



Why Sigmoid?

Function	Range	Differentiable?	Pros	Cons
Step (Heaviside)	$\{0, 1\}$	No	Simple interpretation (hard decision)	Not smooth → no gradient, can't optimize
Linear (identity / scaling)	$(-\infty, \infty)$	Yes	Easy to compute	Not bounded in $[0, 1]$, invalid probability
Tanh	$(-1, 1)$	Yes	Symmetric, useful in neural networks	Output not in $(0, 1)$; needs rescaling
Softmax	$[0, 1]$, sums to 1 across classes	Yes	Best for multi-class problems	Overkill for binary classification
Sigmoid (logistic)	$(0, 1)$	Yes	Perfect probability mapping; convex loss; interpretable log-odds	Can saturate (slow learning at extremes)

Logistic Regression

- Performance measure: Log loss function

$$J(\theta) = -\log L(\theta) = -\sum_{i=1}^m [y^{(i)} \log h_\theta(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_\theta(x^{(i)}))]$$

It is easier to explain the log loss function with one train example case, in which we want to maximize the posterior probability

$$\Pr(y|x) = [h_\theta(x)]^y [(1 - h_\theta(x))^{1-y}] = \begin{cases} h_\theta(x), & \text{when } y = 1 \\ 1 - h_\theta(x), & \text{when } y = 0 \end{cases}$$



$$\log \Pr(y|x) = y \log h_\theta(x) + (1 - y) \log(1 - h_\theta(x))$$

Logistic Regression

- Similarly, next step is to choose the parameters θ that minimize the loss function

$$\theta = \operatorname{argmin} \mathcal{J}(\theta)$$

Where $J(\theta) = -\sum_{i=1}^m [y^{(i)} \log h_\theta(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_\theta(x^{(i)}))]$

- No normal equation (i.e., closed-form solution) for θ
- The cost function is convex and derivable. Gradient Descent is guaranteed to find global maximum

Training Logistic Regression

- The gradient of the log loss function is

$$\frac{\partial}{\partial \theta_j} \mathcal{J}(\theta) = \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

- At each round of gradient descent, θ is updated as following

$$\theta_j = \theta_j - \frac{1}{m} \sum_{i=1}^m \alpha (h_\theta(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

Similar to the linear regression, you can choose different mode of gradient descent and add regularization to the loss function

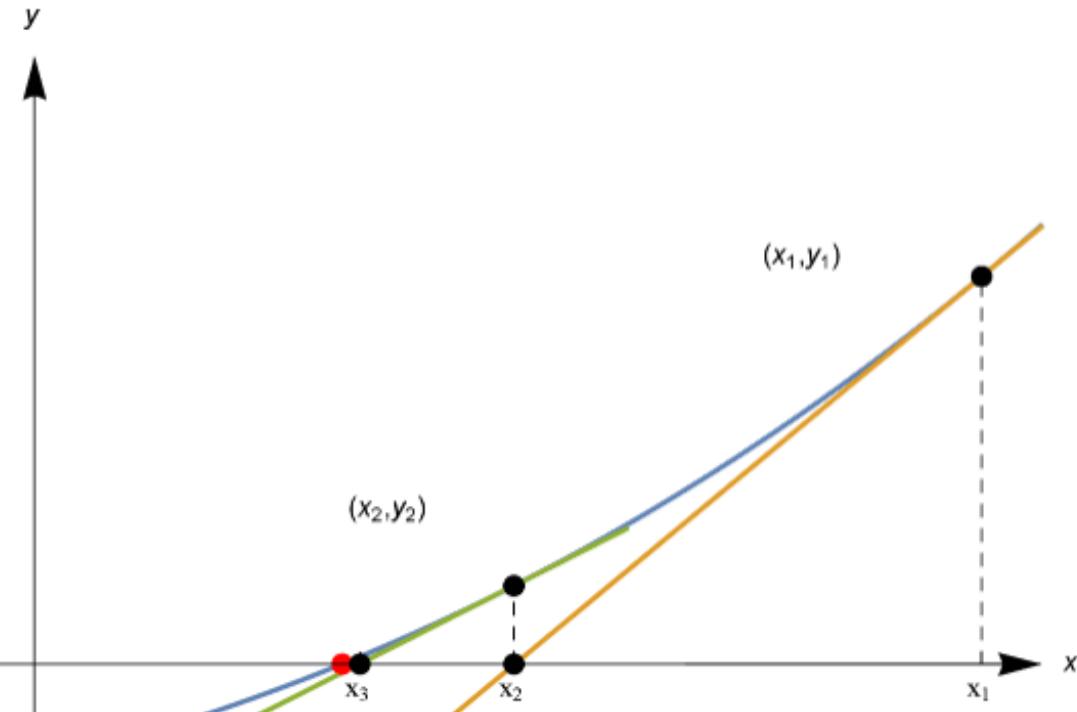
Newton's Method

- Idea: using tangent lines
- Why: efficient, move a much bigger step.
- Process:
 1. Find the tangent.
 2. Find the intersects of tangent line and x-axis
 3. Repeat
- Example:

In (x_1, y_1) , the tangent line:

$$\frac{y - y_1}{x - x_1} = f'(x_1)$$

When we set $y=0$, we have $x_2 = x_1 - \frac{y_1}{f'(x_1)}$ → $x_{n+1} = x_n - \frac{f(x_n)}{f'(x_n)}$



Newton's Method in LR

For LR, we are looking for minimizing the cost function, so actually we set the $f(x) = \mathcal{J}'(\theta)$

Thus, the updating rule:

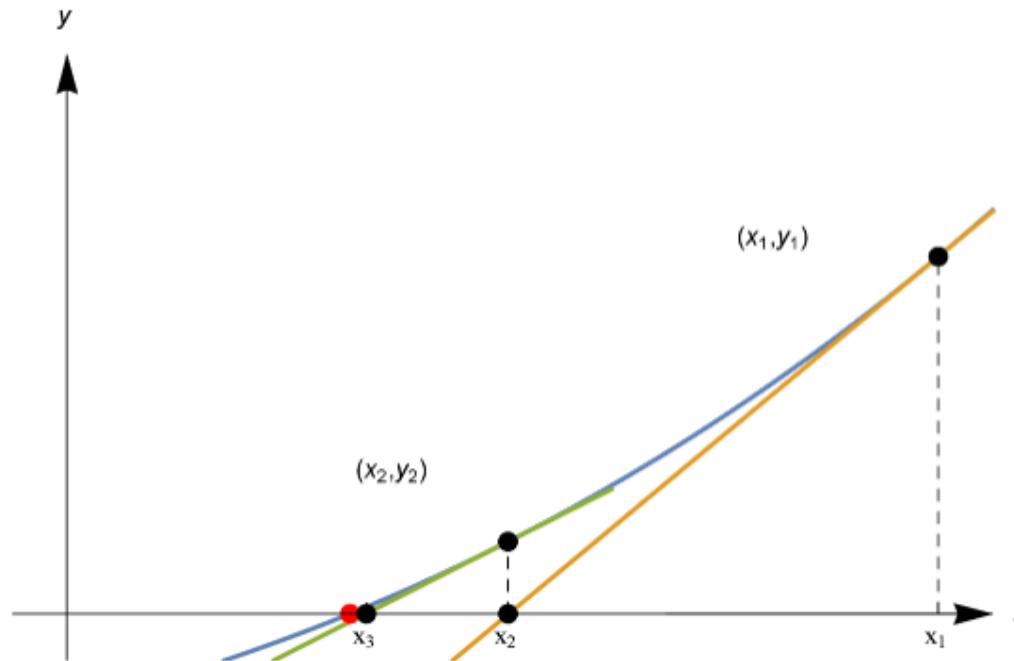
$$\theta := \theta - \frac{\mathcal{J}'(\theta)}{\mathcal{J}''(\theta)}$$

Notice here, θ is a real number which means it is one dimension. It can be generalized to the vector (Newton-Raphson Method):

$$\theta := \theta - H^{-1} \cdot \mathcal{J}'(\theta)$$

Where H is Hessian matrix

$$H_{ij} = \frac{\partial^2 \mathcal{J}(\theta)}{\partial \theta_i \partial \theta_j}, i, j \in [0, n]$$



Imagine x is the parameter axis and y is the corresponding derivatives of cost function

Newton's Method in LR

- Advantages:
 - Not sensitive to the initialization
 - Quadratic convergence rate
- Disadvantages:
 - Much more expensive when we have a ``high-dimensional'' dataset
 - Require the cost function to be twice differentiable
 - The method only works for convergent functions and will be not effective in multiple curvature optimization (neural network)

Multi-class Classification

- “One -vs-Rest” method:

We can use **binary classifier** for multi-class classification with so-called the One-Vs-Rest (OvR) method. Specifically, it uses multiple rounds of binary classification for multi-class classification.

For example, to determine if an object X is a dog, cat or fish, we call a binary classifier $f_i()$ as follows:

```
if  $f_1(X)$  outputs dog  
    return dog;  
else if  $f_2(X)$  outputs cat  
    return cat;  
else return fish
```

Multi-class Classification

One-Vs-One Method:

The One-Vs-One (OvO) method constructs a **binary classifier** for each pair of classes. Therefore, with K classes, we need to construct $K(K-1)/2$ binary classifiers.

The decision at prediction time can be made by **counting the votes** from individual binary classifiers. In case of a tie, it compares the aggregated classification confidence (i.e., the output probability) of individual binary classifiers of each class and the higher one is selected.

The OvO method is slower than OvR. But for some algorithms (e.g., Kernel algorithms) which cannot scale with many training examples, this algorithm can be helpful.

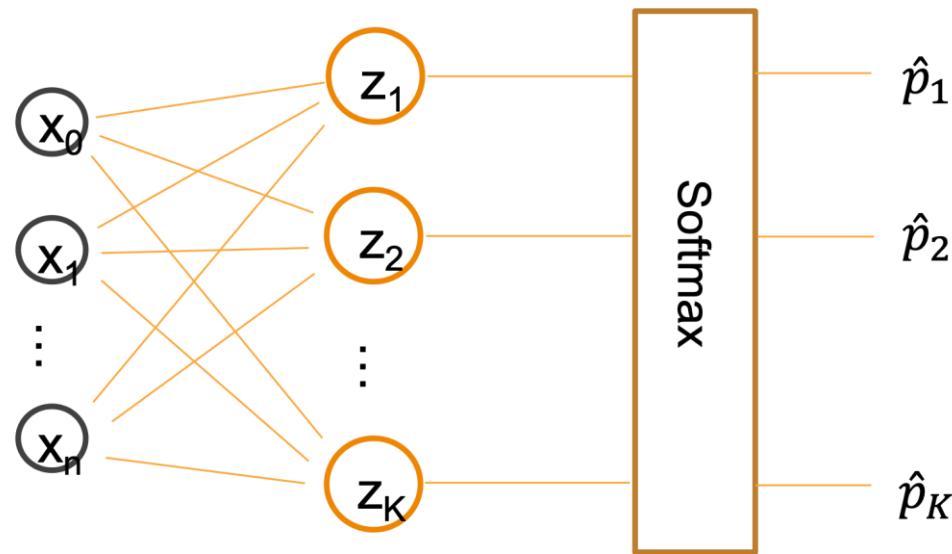
Multi-class Classification

Another approach for multi-class classification is to use the ***multinomial logistic regression***. For each class $1 \leq k \leq K$,

- 1) first compute $z_k(x) = \theta_k^T \cdot x$
- 2) then compute Softmax function:

$$h_k(x) = g(z_k(x))_k = \frac{e^{z_k(x)}}{\sum_{j=1}^K e^{z_j(x)}}$$

where θ_k is the vector of parameters of input features for z_k .



Multi-class Classification

- The performance measure is the cross-entropy cost function

$$\mathcal{J}(\theta) = \frac{1}{n} \sum_{i=1}^n \sum_{k=1}^K y_k^{(i)} \log(h_k(x^{(i)}))$$

where $y_k^{(i)} = \begin{cases} 1, & \text{if } x^{(i)} \text{ is an example of class } k \\ 0, & \text{if } x^{(i)} \text{ is not an example of class } k \end{cases}$

- The gradient descent methods can be used to train the multinomial logistic regression model. And the gradients can be computed as:

$$\frac{\partial}{\partial \theta_k} \mathcal{J}(\theta) = \frac{1}{n} \sum_{i=1}^n (h_k(x^{(i)}) - y_k^{(i)}) x^{(i)}$$

Multi-class Classification

Other Approaches: Error-Correcting Output Codes

The Error-Correcting Output Codes (ECOC) method encodes **K classes** into **N bit vectors**. Each class is represented as a bit in each bit vector. ECOC trains **N binary classifiers**, each splitting one group of classes from another (using the column bit vectors below). At prediction time, the N binary classifiers are called, the outputs of them yielding an N -bit vector. A class with the **closest Euclidean distance** to the **N -bit vector** is selected. To reduce the classification error, error correcting codes are used when generating the “code book”

Class	Code Word														
	f_0	f_1	f_2	f_3	f_4	f_5	f_6	f_7	f_8	f_9	f_{10}	f_{11}	f_{12}	f_{13}	f_{14}
0	1	1	0	0	0	0	1	0	1	0	0	1	1	0	1
1	0	0	1	1	1	1	0	1	0	1	1	0	0	1	0
2	1	0	0	1	0	0	0	1	1	1	1	0	1	0	1
3	0	0	1	1	0	1	1	1	0	0	0	0	1	0	1
4	1	1	1	0	1	0	1	1	0	0	1	0	0	0	1
5	0	1	0	0	1	1	0	1	1	1	0	0	0	0	1
6	1	0	1	1	1	0	0	0	0	1	0	1	0	0	1
7	0	0	0	1	1	1	0	1	0	1	1	0	0	0	1
8	1	1	0	1	0	1	1	0	0	1	0	0	0	1	1
9	0	1	1	1	0	0	0	0	1	0	1	0	0	1	1

Classifier Evaluation

Classification measures

- Accuracy is only one measure (error = 1-accuracy).

Accuracy is not always suitable

- In text mining, we may only be interested in the documents of a particular topic, which are only a small portion of a big document collection.
- In classification involving skewed or highly imbalanced data, e.g., network intrusion and financial fraud detections, **we are interested only in the minority class.**
 - High accuracy does not mean any intrusion is detected.
 - E.g., 1% intrusion. Achieve 99% accuracy by doing nothing.
- The class of interest is commonly called the **positive class**, and the rest **negative classes**.

Performance Measures of Classifiers

Sometimes we need both a high accuracy rate and a low error rate. **Confusion matrix** is an important tool to measure both accuracy rates and error rates.

	Classified Positive	Classified Negative
Actual Positive	TP	FN
Actual Negative	FP	TN

True Positive (TP): the number of true positive examples

True Negative (TN): the number of true negative examples

False Positive (FP): the number of false positive examples (type I error)

False Negative (FN): the number of false negative examples (type II error)

Performance Measures of Classifiers

We are often interested in two measures on relevance:

- **precision** $p = \frac{TP}{TP+FP}$, is the number of true positive divided by all classified as positive.
- **recall (a.k.a., sensitivity)** $r = \frac{TP}{TP+FN}$, is the number of true positive divided by all positive.
- $F_1 = \frac{2pr}{p+r}$, is a measure of both p and r . F_1 -score is large only if both p and r are large.

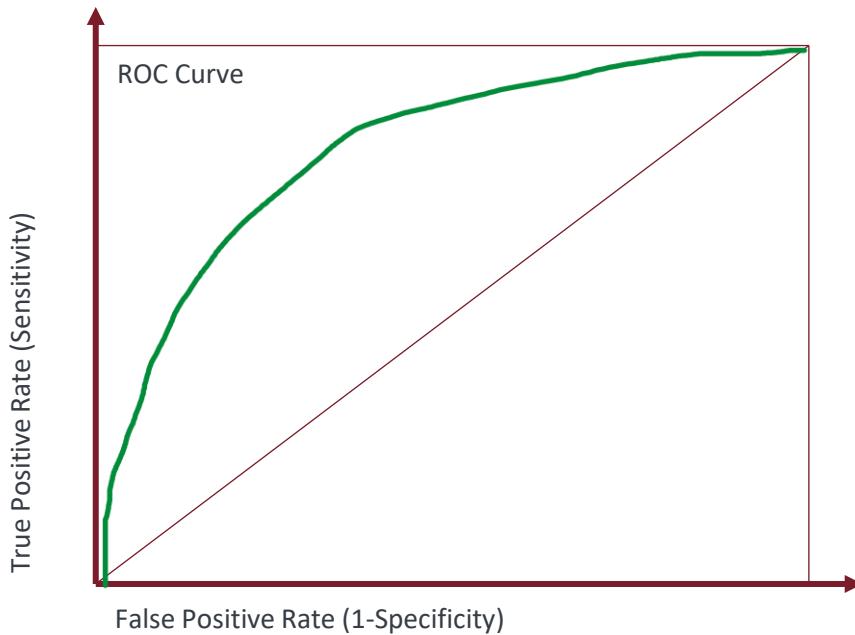
	Classified Positive	Classified Negative
Actual Positive	TP	FN
Actual Negative	FP	TN

“**Precision** can be seen as a measure of quality and **recall** as a measure of quantity.

- Higher precision means that an algorithm returns more relevant results than irrelevant ones,
- high recall means that an algorithm returns most of the relevant results.” (Wikipedia)

Performance Measures of Classifiers

Another common measure is the **Receiver Operating Characteristic** (ROC) curve, which shows the tradeoff between **True Positive Rate** (i.e., $sensitivity = \frac{TP}{TP+FN}$) and **False Positive Rate** (i.e., $1 - specificity = \frac{FP}{FP+TN}$).



- The closer to the top-left corner of the ROC space, the better the classifier is.
- The closer to the 45-degree diagonal of the ROC space, the worse the classifier is.
- The area under the curve (AUC) is usually used to measure above properties.
- ROC/AUC are usually used to measure **how well a binary classifier distinguish the two classes**.



STEVENS
INSTITUTE OF TECHNOLOGY
1870

THANK YOU

Stevens Institute of Technology
1 Castle Point Terrace, Hoboken, NJ 07030

