

AAI/CPE/EE 595-WS1

Applied Machine Learning

Lecture 5-2: Ensemble Learning

Speaker: Jiarui Li
jli148@stevens.edu

Classification

Task: find a function/hypothesis h that maps the input space \mathcal{X} to a set of classes $\mathcal{Y} = \{y_1, \dots, y_n\}$

- Most popular single classifiers:
 - Logistic Regression
 - Decision Tree
 - Bayes Classifier
 - K-Nearest Neighbor
 - Neural Network
 - Support Vector Machine

Drawback of Single Classifier

The “best” classifier not necessarily the ideal choice

- Problems:
 - Which one is the best?
 - Maybe more than one classifiers meet the criteria (e.g. same training accuracy), especially in the following situations:
 - Without sufficient training data
 - The learning algorithm leads to different local optima easily
 - Potentially valuable information may be lost by discarding the results of less-successful classifiers
 - E.g., the discarded classifiers may correctly classify some samples
 - The trained classifier may not be complex enough to handle the problem

Classifier Ensemble

- Combining a number of trained classifiers lead to a better performance than any single one
 - Errors can be complemented by other correct classifications
 - Different classifiers have different knowledge regarding the problem
- To decompose a complex problem into sub-problems for which the solutions obtained are simpler to understand, implement, manage and update

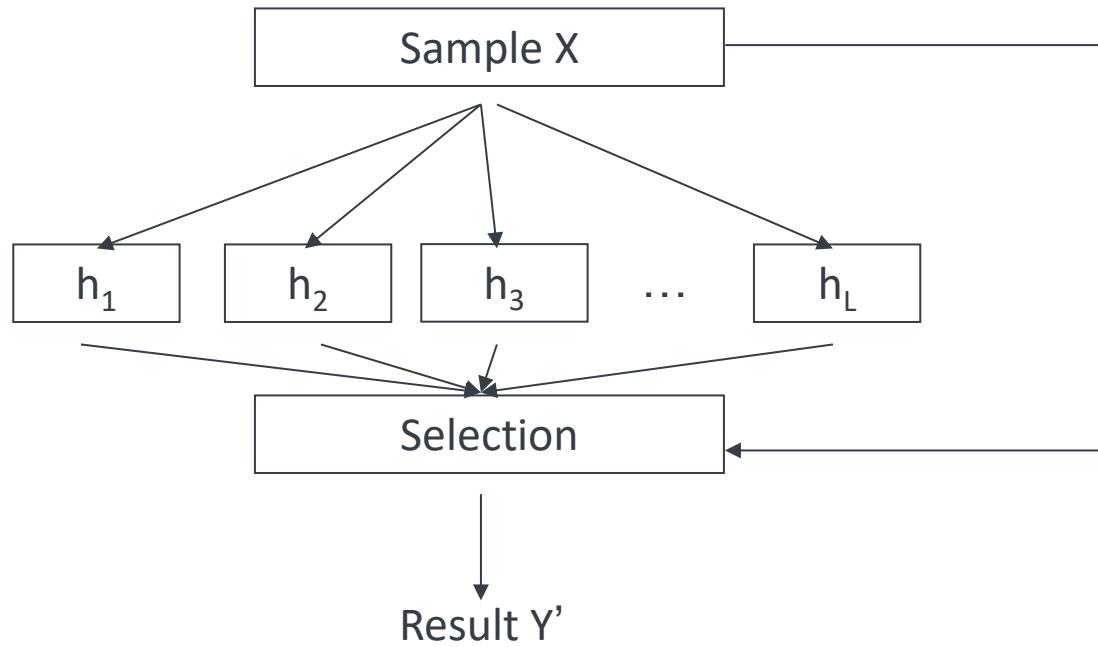
Classifier Ensemble

- Classifier ensemble consists of
 - a set of individual classifiers
 - a fusion/selection method:
to combine/select individual classifier outputs to give a final decision
- Types of ensemble:
 - Classifier Selection
 - Classifier Fusion

Classifier Selection

- The input sample space is partitioned into smaller areas and each classifier learns the sample in each area
- For each sample, identify a single classifier which is most likely to produce the correct classification label
- Only the output of the selected classifier is taken as a final decision
- It is similar to the “Divide and Conquer” approach

Classifier Selection



Classifier ensemble MUST be better than Single?

- In all cases? NO!
- But in many cases, a ensemble learning can have a better performance than a single classifier

Affecting Factors

Three factors affecting the ensemble model accuracy:

① Accuracy of individual classifiers

How good are the individual classifiers?

- Training Dataset (sample and feature)
- Learning Model (types of classifier)
- Model's Parameters (e.g. the number of neurons in NN)
-

② Fusion Methods

How to combine classifiers?

③ Diversity among classifiers

What are the differences among different classifiers?

Fusion Method

- A method to combine individual classifier outputs to reach the final decision for the classifier ensemble learning
- Since different fusion methods may have different final outputs for the same individual classifier outputs, **MCS** error is affected by its fusion method
- For Example

Sample x	Class1	Class2
Classifier 1	0.2	0.8
Classifier 2	0.7	0.3
Classifier 3	0.7	0.3

average Class 1: 0.53
 Class 2: 0.47

max Class 1: 0.7
 Class 2: 0.8

- Two popular fusion methods based on individual classifier outputs types:
 - Soft Output: soft type rules
 - Class Label: hard type rules

Soft Output

- For each sample, the classifier outputs a value representing the **confidence** of this sample belonging to each class
- The output of a classifier i is a c -dimensional vector $[d_{i,1}, d_{i,2}, \dots, d_{i,c}]^T$, where c is number of classes
- It is called decision profile

	y_1	y_2
Classifier 1	0.5	0.5
Classifier 2	0.2	0.8
Classifier 3	0.6	0.4

Class Label

- A classifier outputs only one class label

Classifier	output
Classifier 1	y_1
Classifier 2	y_2
Classifier 3	y_1

Fusion Method: SSF

- SSF stands for **Simple Summary Function**. It uses simple summary of different classifiers to output the final decision. E.g., average, maximum, minimum or product. These fusion methods calculate the support $\mu_j(x)$ for class j

AVR:

$$\mu_j(x) = \frac{1}{L} \sum_{i=1}^L d_{i,j}(x)$$

MIN:

$$\mu_j(x) = \min_{i,j}$$

MAX:

$$\mu_j(x) = \max_{i = 1, \dots, L_{i,j}}$$

PRO:

$$\mu_j(x) = \prod_{i=1}^L d_{i,j}(x)$$

Simple Summary Function(SSF)

- An example:
 - Average
 - **Class 1: 0.5**
 - Class 2: 0.5
 - Minimum
 - Class 1: 0.2
 - **Class 2: 0.3**
 - Maximum
 - Class 1: 0.7
 - **Class 2: 0.8**
 - Product
 - Class 1: 0.084
 - **Class 2: 0.096**

Sample x	Class1	Class2
Classifier 1	0.2	0.8
Classifier 2	0.6	0.4
Classifier 3	0.7	0.3

Fusion Method: Weighted Average

- The Weighted Average (WAVR)

The support for class j :

$$\mu_j(x) = \frac{1}{L} \sum_{i=1}^L w_i d_{i,j}(x)$$

where $\sum_{i=1}^L w_i = 1$

Usually, the weight is calculated using the accuracy of the individual classifier

Weighted Average

- An example:
 - Case 1:
 - Assume that the weight of
 - Classifier 1: 0.7
 - Classifier 2: 0.2
 - Classifier 3: 0.1
 - Class 1: 0.33
 - Class 2: 0.67
 - Case 2:
 - Assume that the weight of
 - Classifier 1: 0.2
 - Classifier 2: 0.3
 - Classifier 3: 0.5
 - Class 1: 0.57
 - Class 2: 0.43

x	Class1	Class2
Classifier 1	0.2	0.8
Classifier 2	0.6	0.4
Classifier 3	0.7	0.3

Fusion Method for Class Label

- Popular methods :
 - Majority Vote (MV)
 - Weighted Majority Vote (WMV)
 - Naïve Bayes (NB)

Majority Vote

- This method may be the oldest and the most well-known strategy for decision making
- Assume that the label outputs of the classifiers are given as c -dimensional binary vectors
 $[d_{i,1}, d_{i,2}, \dots, d_{i,c}]^T$ in $\{0,1\}^c$
- $i = 1, \dots, L$ where $d_{i,j} = 1$ if label x in class j , and 0 otherwise
- The majority vote results in an ensemble decision for class k if

$$\sum_{i=1}^{L \Sigma} d_{i,k} = \max_{j=1}^c$$

- An example:
 - Class 1: 2 votes
 - **Class 2: 3 votes**

Sample x	Result
Classifier 1	1
Classifier 2	2
Classifier 3	1
Classifier 4	2
Classifier 5	2

Weighted Majority Vote

- Similar to majority vote method but the influence of each vote to the final decision is not the same
- The weighted majority vote results in an ensemble decision for class k if

$$\sum_{i=1}^{L \sum_{i=1}^L w_i d_{i,j}} w_i d_{i,k} = \max_{j=1}^c$$

- An example:
 - Case 1:
 - Assume that the weight of
 - Classifier 1: 0.1
 - Classifier 2: 0.2
 - Classifier 3: 0.2
 - Classifier 4: 0.3
 - Classifier 5: 0.2
 - Class 1: 0.3
 - Class 2: 0.7
 - Case 2:
 - Assume that the weight of
 - Classifier 1: 0.4
 - Classifier 2: 0.2
 - Classifier 3: 0.2
 - Classifier 4: 0.1
 - Classifier 5: 0.1
 - Class 1: 0.6
 - Class 2: 0.4

Sample x	Result
Classifier 1	Class 1
Classifier 2	Class 2
Classifier 3	Class 1
Classifier 4	Class 2
Classifier 5	Class 2

Naïve Bayes

- The term “naïve” is used since this method relies on the assumption that the classifiers are mutually independent but this situation does not occur normally

$$\mu_j(x) \propto \prod_{i=1}^L \hat{P}(\omega_j | d_{i,j}(x) = 1)$$

- $\hat{P}(\omega_j | d_{i,j}(x) = 1)$ is learned from training samples

Naïve Bayes

- An example:
 - The outputs of individual classifiers are **1 2 1**.

$$\begin{aligned}\hat{P}(\omega_1|d_{1,1}(x) = 1) &= \frac{40}{70} & \hat{P}(\omega_2|d_{1,1}(x) = 1) &= \frac{30}{70} \\ \hat{P}(\omega_1|d_{2,2}(x) = 1) &= \frac{30}{60} & \hat{P}(\omega_2|d_{2,2}(x) = 1) &= \frac{30}{60} \\ \hat{P}(\omega_1|d_{3,1}(x) = 1) &= \frac{50}{90} & \hat{P}(\omega_2|d_{3,1}(x) = 1) &= \frac{40}{90}\end{aligned}$$

↓

Class 1: 0.16
Class 2: 0.10

```
graph TD; C1["Class 1: 0.16"] --- B1[""]; C2["Class 2: 0.10"] --- B2[""]; B1 --- G1[""]; B2 --- G2[""]; G1 --- S1[""]; G2 --- S2[""]; S1 --- C1; S2 --- C2;
```

		Classifier1 Decision	
		Class1	Class2
True	Class1	40	10
	Class2	30	20

		Classifier2 Decision	
		Class1	Class2
True	Class1	20	30
	Class2	20	30

		Classifier3 Decision	
		Class1	Class2
True	Class1	50	0
	Class2	40	10

Major Method

- Bootstrap aggregating (Bagging)
- Boosting (AdaBoost)
- Stacked generalization
- Random subspace method
- Mixture of experts

Bagging & Pasting

- Use the same training algorithm, each time training with a random subset of samples of a training set.
 - If random sampling with repeat, it is called *bagging*
 - If random sampling without repeat, it is called *pasting*
- Several classifiers will be trained from the same original dataset.
- Make decision by aggregating multiple classifiers (e.g., majority vote)
- This help decrease the variance (hence the error) in the results in *unstable learners* (e.g., decision trees and neural networks) whose output can change dramatically when the training data is slightly changed.
- The Random Forest algorithm uses the Bagging (or sometimes pasting) method.

Bagging

Input: sequence of m examples $\langle(x_1, y_1), \dots, (x_m, y_m)\rangle$ with labels $y_i \in Y = \{1, \dots, k\}$

weak learning algorithm **WeakLearn**

integer T specifying number of iterations

- Do $t=1, 2, \dots, T$
 - Obtain bootstrap sample S_t by randomly drawing \underline{m} instances, with replacement, from the original training set;
 - Call the WeakLearn based on the S_t to build a hypothesis;
- Using majority voting to combine these T individual hypothesis

Boosting

- Like bagging, but trying to correct previous classifiers at each round
 - When sampling training set, higher priority is given to examples that caused errors to previous hypothesis.
 - The trained hypothesis will automatically overcome the "difficult" training examples.
 - Final decision is made by weighted majority vote.
- The AdaBoost algorithm uses this approach.
- Another example of such is the Gradient Boosting algorithm.

Adaptive Boosting - AdaBoost

- Similar to bagging method, each individual classifier is also trained using a different training set
 - But the training set is selected based on the error of the trained hypothesis: more weights given to the “difficult” examples!
 - Finally, Weight Majority Voting (WMV) will be used as a fusion method
-
- Key Insights
 - Instead of sampling (as in bagging), re-weight examples!
 - Examples are given weights. At each iteration, a new hypothesis is learned (**weak learner**) and the examples are reweighted to focus the system on examples that the most recently learned classifier got wrong.
 - Final classification based on **weighted vote of weak classifiers**

AdaBoost.M1

Algorithm AdaBoost.M1

Input: sequence of m examples $\langle(x_1, y_1), \dots, (x_m, y_m)\rangle$ with labels $y_i \in Y = \{1, \dots, k\}$

weak learning algorithm **WeakLearn**

integer T specifying number of iterations

Initialize $D_1(i) = 1/m$ for all i .

Do for $t = 1, 2, \dots, T$

1. Call **WeakLearn**, providing it with the distribution D_t .

2. Get back a hypothesis $h_t : X \rightarrow Y$.

3. Calculate the error of h_t : $\epsilon_t = \sum_{i:h_t(x_i) \neq y_i} D_t(i)$. If $\epsilon_t > 1/2$, then set $T = t - 1$ and abort loop.

4. Set $\beta_t = \epsilon_t / (1 - \epsilon_t)$.

5. Update distribution D_t : $D_{t+1}(i) = \frac{D_t(i)}{Z_t} \times \begin{cases} \beta_t & \text{if } h_t(x_i) = y_i \\ 1 & \text{otherwise} \end{cases}$

where Z_t is a normalization constant (chosen so that D_{t+1} will be a distribution).

Output the final hypothesis: $h_{fin}(x) = \arg \max_{y \in Y} \sum_{t:h_t(x)=y} \log \frac{1}{\beta_t}$.

AdaBoost.M2

Algorithm AdaBoost.M2

Input: sequence of m examples $\langle (x_1, y_1), \dots, (x_m, y_m) \rangle$ with labels $y_i \in Y = \{1, \dots, k\}$

weak learning algorithm **WeakLearn**

integer T specifying number of iterations

Let $B = \{(i, y) : i \in \{1, \dots, m\}, y \neq y_i\}$

Initialize $D_1(i, y) = 1/|B|$ for $(i, y) \in B$.

Do for $t = 1, 2, \dots, T$

1. Call **WeakLearn**, providing it with mislabel distribution D_t .
2. Get back a hypothesis $h_t : X \times Y \rightarrow [0, 1]$.
3. Calculate the pseudo-loss of h_t : $\epsilon_t = \frac{1}{2} \sum_{(i,y) \in B} D_t(i, y)(1 - h_t(x_i, y_i) + h_t(x_i, y))$.
4. Set $\beta_t = \epsilon_t / (1 - \epsilon_t)$.
5. Update D_t : $D_{t+1}(i, y) = \frac{D_t(i, y)}{Z_t} \cdot \beta_t^{(1/2)(1+h_t(x_i,y_i)-h_t(x_i,y))}$
where Z_t is a normalization constant (chosen so that D_{t+1} will be a distribution).

Output the hypothesis: $h_{fin}(x) = \arg \max_{y \in Y} \sum_{t=1}^T \left(\log \frac{1}{\beta_t} \right) h_t(x, y)$.

Stacking

- Previous ensemble methods use relatively “trivial” functions (e.g., majority voting) to aggregate results of individual classifiers.
- Differently, the Stacking method **train a separate model** to perform the aggregation.
 - Two-layer model: the first layer are the individual classifiers; the second layer is the aggregation model.
 - Training set is divided into two subsets: one is to train the individual classifiers, the other to train the aggregation model (a.k.a., the *hold-out set*).
 - First sufficiently train individual classifiers.
 - Then use the hold-out set to train the aggregation model.
 - Outputs of individual classifiers are inputs to the aggregation model.

Random Subspace Method

- To choose a group of features for each classifier
- Each individual classifier is then trained on the randomly selected group of features
- Their results are combined by a fusion rule.

- Sample: Random forests .

Random Forest

A random forest is a collection of CART-like trees following specific rules for

- Tree growing
- Tree combination
- Self-testing
- Post-processing

Random Forest

Tree Growing/Split

- Binary partitioning
- Each tree is grown at least partially *at random*
 - growing each tree on a different random subsample of the training data
 - splitting at any node from the eligible random subset

Prediction Mechanism

- Grow many trees.
- Each tree casts a vote at its terminal nodes. For a binary target the vote will be YES or NO
- Count up the YES votes. The percent YES votes received is the predicted probability



STEVENS
INSTITUTE OF TECHNOLOGY
1870

THANK YOU

Stevens Institute of Technology
1 Castle Point Terrace, Hoboken, NJ 07030

